

# Docker - Introduction and Basics

## Table of Contents

1. What is Docker?
  2. Why We Need Docker?
  3. What is a Container?
  4. Before Docker - The Traditional Problem
  5. History of Docker
  6. What is DevOps?
- 

## What is Docker?

**Docker** is an open-source platform that automates the deployment, scaling, and management of applications using containerization technology. It packages applications and their dependencies into standardized units called containers.

### Key Features of Docker:

- **Portability:** Run anywhere - development, testing, production
- **Lightweight:** Shares host OS kernel, uses minimal resources
- **Fast:** Containers start in seconds
- **Isolated:** Each container runs independently
- **Scalable:** Easy to scale up or down

### Real-World Analogy:

Think of Docker like shipping containers in logistics:  
- Standard size and shape  
- Can be loaded on any ship, truck, or train  
- Contents are isolated and protected  
- Easy to track and manage

### Example:

```
# Docker allows you to do this:  
docker run -d -p 80:80 nginx  
  
# And boom! You have a web server running in seconds
```

---

## Why We Need Docker?

### Problem 1: “It Works on My Machine” Syndrome

#### Traditional Scenario:

Developer: "The app works perfectly on my laptop!"  
QA Team: "It's broken on our test server"

**Operations:** "It crashes in production"

**Root Causes:** - Different OS versions - Missing dependencies - Different library versions - Configuration differences - Environment variables mismatch

### Problem 2: Dependency Hell

**Example Without Docker:**

Application A requires:

- Python 3.8
- Node.js 14
- PostgreSQL 12
- Redis 5.0

Application B requires:

- Python 3.11
- Node.js 18
- PostgreSQL 15
- Redis 7.0

How do you run both on the same server?

### Problem 3: Resource Wastage with Virtual Machines

**Traditional VM Approach:**

Server (64GB RAM, 16 CPU cores)

- VM1 (16GB RAM, 4 cores) - Runs App A
- VM2 (16GB RAM, 4 cores) - Runs App B
- VM3 (16GB RAM, 4 cores) - Runs App C
- VM4 (16GB RAM, 4 cores) - Runs App D

Each VM runs a full OS = Lots of wasted resources!

**Docker Solutions:**

#### Solution 1: Consistent Environment

```
# Define your environment once
FROM node:18
WORKDIR /app
COPY package.json .
RUN npm install
COPY . .
CMD ["npm", "start"]

# Now it works EVERYWHERE!
```

## Solution 2: Isolated Dependencies

```
# Run multiple versions without conflict
docker run -d --name app-a python:3.8
docker run -d --name app-b python:3.11

# Each container has its own environment
```

## Solution 3: Efficient Resource Usage

```
Server (64GB RAM, 16 CPU cores)
  Container 1 (2GB RAM) - App A
  Container 2 (1GB RAM) - App B
  Container 3 (3GB RAM) - App C
  Container 4 (2GB RAM) - App D
  ...many more containers
```

All share the same OS kernel = More efficient!

### Benefits Summary:

Benefit	Description	Example
<b>Consistency</b>	Same environment everywhere	Dev = Test = Production
<b>Isolation</b>	Apps don't interfere	Multiple Node.js versions
<b>Portability</b>	Run on any Docker host	AWS, Azure, Local
<b>Speed</b>	Fast startup and deployment	Start in milliseconds
<b>Scalability</b>	Easy to scale horizontally	Run 100 containers easily
<b>Version Control</b>	Track image versions	Roll back easily

---

## What is a Container?

### Definition

A **container** is a lightweight, standalone, executable package that includes everything needed to run a piece of software: - Code - Runtime - System tools - System libraries - Settings

### Container Characteristics:

#### 1. Lightweight

Virtual Machine: 1-2 GB+  
Container: 50-200 MB

Why? Containers share the host OS kernel!

## 2. Isolated

```
# Each container has its own:  
- Filesystem  
- Process space  
- Network interface  
- Users and groups  
  
# Example: Two containers, same port, no conflict  
docker run -d -p 8080:80 nginx # Container 1  
docker run -d -p 8081:80 nginx # Container 2
```

## 3. Portable

```
# Build on Mac  
docker build -t myapp .  
  
# Run on Linux server  
docker run myapp  
  
# Run on Windows  
docker run myapp  
  
# Same result everywhere!
```

### Container vs Process

#### Regular Process:

Your App → Shares everything with host  
(filesystem, network, users)

#### Container:

Your App → Isolated environment  
Own filesystem  
Own network  
Own processes  
Controlled resources

### Real-World Example:

**Scenario:** Running a Node.js application

**Without Container:**

```
# Install Node.js globally
sudo apt install nodejs

# Install dependencies
npm install

# Run app
node app.js

# Problem: What if another app needs different Node.js version?
```

**With Container:**

```
# Dockerfile
FROM node:18-alpine
WORKDIR /app
COPY package*.json ./
RUN npm install
COPY .
EXPOSE 3000
CMD ["node", "app.js"]

# Build
docker build -t my-node-app .

# Run
docker run -p 3000:3000 my-node-app

# Benefits:
# Specific Node.js version
# All dependencies included
# Doesn't affect host system
# Can run multiple versions simultaneously
```

**Container Lifecycle:**

1. CREATE → Container is created from image  
docker create nginx
2. START → Container starts running  
docker start container\_id
3. RUNNING → Container is actively running  
docker ps
4. STOP → Container stops gracefully

```
        docker stop container_id

5. REMOVE      → Container is deleted
                docker rm container_id
```

**Example:**

```
# Complete lifecycle example
docker create --name my-nginx nginx
# Output: container_id_12345

docker start my-nginx
# Output: my-nginx

docker ps
# Shows running container

docker stop my-nginx
# Gracefully stops container

docker rm my-nginx
# Removes container completely
```

---

## Before Docker - The Traditional Problem

The Dark Ages of Application Deployment (Pre-2013)

Problem 1: Manual Server Configuration Traditional Setup:

```
# On Server 1
ssh admin@server1.company.com

# Install dependencies manually
sudo apt update
sudo apt install python3.8
sudo apt install postgresql
sudo apt install redis-server
sudo apt install nginx

# Configure each service
sudo nano /etc/postgresql/postgresql.conf
sudo nano /etc/nginx/nginx.conf
sudo nano /etc/redis/redis.conf

# Deploy application
git clone https://github.com/company/app.git
```

```
cd app
pip install -r requirements.txt
python manage.py migrate
python manage.py runserver

# Problem: Need to repeat this on 50 servers?
```

### Problem 2: Inconsistent Environments The Conversation:

Developer (Monday 9 AM):  
"I built this awesome feature over the weekend!"  
Runs: python manage.py runserver  
Works perfectly!

QA Engineer (Monday 2 PM):  
"The feature is broken, can't even start the app!"  
Runs: python manage.py runserver  
Error: Module 'numpy' not found

Operations (Tuesday 10 AM):  
"Production is down! Your feature crashed everything!"  
Runs: python manage.py runserver  
Error: Python version mismatch

Developer:  
"But... it works on my machine!"

### Why It Happened:

Developer's Laptop:  
- Ubuntu 22.04  
- Python 3.10  
- PostgreSQL 14  
- numpy 1.24.0

QA Server:  
- Ubuntu 20.04  
- Python 3.8  
- PostgreSQL 12  
- numpy not installed

Production Server:  
- Ubuntu 18.04  
- Python 3.6  
- PostgreSQL 10  
- numpy 1.19.0

### Problem 3: Dependency Conflicts Scenario:

```
# Company runs multiple projects on one server

# Project A requirements:
Django==3.2
PostgreSQL==12

# Project B requirements:
Django==4.2
PostgreSQL==15

# Try to install both:
pip install Django==3.2 # For Project A
pip install Django==4.2 # For Project B

# Result: Only one version can exist!
# Project A breaks when Project B is deployed
```

### Problem 4: Scaling Nightmare Traditional Scaling:

Step 1: Buy new server (\$\$\$)  
Step 2: Wait for delivery (days/weeks)  
Step 3: Install OS (2 hours)  
Step 4: Install dependencies (3 hours)  
Step 5: Configure services (2 hours)  
Step 6: Deploy application (1 hour)  
Step 7: Test everything (2 hours)  
Step 8: Debug issues (??? hours)

Total Time: Days or weeks  
Total Cost: Thousands of dollars  
Success Rate: 60%

### Problem 5: Server Snowflakes Each server becomes unique over time:

Server 1:  
- Deployed in 2018  
- Manually patched 47 times  
- Configuration files modified by 5 different admins  
- Nobody knows what's different anymore

Server 2:  
- Deployed in 2019  
- Slightly different setup  
- Different versions of libraries  
- Works differently than Server 1

Result: "Snowflake servers" - each one is unique and fragile

### Enter Virtualization (First Solution Attempt)

Virtual Machines to the Rescue... Sort of:

Physical Server

Hypervisor (VMware, VirtualBox)  
VM1 (full OS + App A) - 8GB RAM  
VM2 (full OS + App B) - 8GB RAM  
VM3 (full OS + App C) - 8GB RAM

Problems:

Each VM needs full OS (GBs of disk space)  
Slow to start (minutes)  
Resource intensive  
Still requires manual configuration per VM  
But at least isolated!

### Why Docker Changed Everything

With Docker (2013 onwards):

Physical Server

Docker Engine  
Container A (App only) - 200MB  
Container B (App only) - 150MB  
Container C (App only) - 180MB

Benefits:

Lightweight (MBs not GBs)  
Fast to start (seconds)  
Efficient resource usage  
Consistent across all environments  
Easy to scale  
Infrastructure as Code

Docker Workflow:

```
# Day 1: Developer creates Dockerfile
cat > Dockerfile << EOF
FROM python:3.10
WORKDIR /app
COPY requirements.txt .
RUN pip install -r requirements.txt
COPY . .
CMD ["python", "manage.py", "runserver"]
```

```
EOF

# Day 2: Build image
docker build -t myapp:1.0 .

# Day 3: Run anywhere
docker run -p 8000:8000 myapp:1.0

# Day 4: Scale to 100 servers
for i in {1..100}; do
    docker run -d -p 800$i:8000 myapp:1.0
done

# It just works! Everywhere! Every time!
```

---

## History of Docker

### Timeline of Containerization

#### 1979 - Unix V7

- chroot command introduced
  - Change root directory for a process
  - First step toward isolation

#### 2000 - FreeBSD Jails

- More advanced isolation
  - Separate filesystem
  - Network stack
  - Process space

#### 2001 - Linux VServer

- Operating system-level virtualization
  - Multiple Linux instances on one server

#### 2004 - Solaris Containers

- Sun Microsystems introduces zones
  - Similar to modern containers

#### 2005 - OpenVZ

- Operating system-level virtualization for Linux
  - Patched Linux kernel

## **2006 - Process Containers (cgroups)**

Google introduces cgroups (Control Groups)

- Limit and isolate resource usage
- Foundation for modern containers

## **2008 - LXC (Linux Containers)**

First complete Linux container manager

- Uses cgroups and namespaces
- Template-based

## **2013 - Docker is Born!**

Solomon Hykes demos Docker at PyCon

- Makes containers easy to use
- Introduces Docker Hub
- Standardizes container format

March 2013: Docker 0.1 released

October 2013: Docker gets \$15M funding

## **2014 - The Docker Revolution**

- Docker 1.0 released (production-ready)
- Microsoft announces Windows Server containers
- Google announces Kubernetes
- AWS announces EC2 Container Service

## **2015 - Open Container Initiative (OCI)**

- Docker donates container format to OCI
- Standardizes container specification
- Industry-wide adoption

## **2016-2020 - Container Ecosystem Exploses**

- Docker Swarm for orchestration
- Kubernetes becomes dominant orchestrator
- Docker Desktop for Mac/Windows
- Enterprise adoption soars

## **2021-Present - Cloud Native Era**

- Containers everywhere
- Serverless containers
- Edge computing with containers
- Docker becomes industry standard

## Docker's Impact in Numbers

**2013 (Launch Year):** - Few thousand users - Small community

**2015:** - 1 billion+ container pulls - 100,000+ Dockerized applications

**2020:** - 13+ million developers - 5.6 million repositories on Docker Hub

**2025:** - Industry standard for containerization - Used by 90%+ of Fortune 500 companies

---

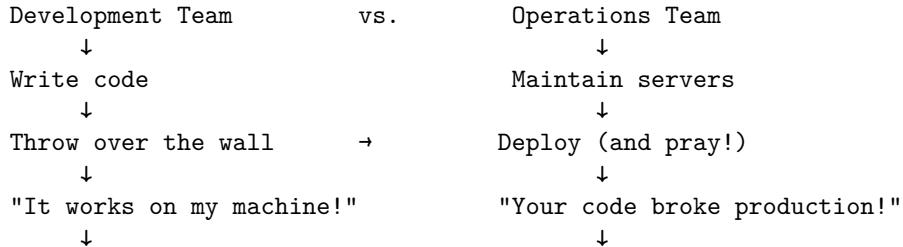
## What is DevOps?

### Definition

**DevOps** is a culture, methodology, and set of practices that combines software Development and IT Operations to shorten the development lifecycle and deliver high-quality software continuously.

### Before DevOps (The Waterfall Era)

#### Traditional Approach:



Timeline: Months or years per release

Quality: Lots of bugs in production

Communication: Hostile

Blame game: Epic

#### The Wall of Confusion:

Developers say:		Operations say:
"Deploy faster!"		"Deploy safely!"
"I need more features!"		"I need stability!"
"It worked in dev!"		"You didn't test properly!"
"Not my problem!"		"Fix your code!"

## The DevOps Philosophy

### Core Principles:

## 1. Culture - Break Down Silos

Before:

Dev Team → Operations Team (enemies)

After:

DevOps Team → Shared responsibility  
- Developers on-call for their code  
- Operations involved in development  
- Shared goals and metrics

## 2. Automation - Automate Everything

```
# Manual deployment (old way)
ssh server1.com
cd /var/www/app
git pull
npm install
systemctl restart app
# Repeat for 50 servers

# Automated deployment (DevOps way)
git push origin main
# CI/CD pipeline automatically:
# - Tests code
# - Builds containers
# - Deploys to all servers
# - Monitors health
# - Rolls back if issues
# Done in minutes!
```

## 3. Measurement - Monitor Everything

Track:

- Deployment frequency
- Lead time for changes
- Mean time to recovery
- Change failure rate
- Application performance
- User experience

## 4. Sharing - Knowledge and Tools

Share:

- Code repositories (Git)
- Documentation (Wiki)
- Runbooks

- Monitoring dashboards
- On-call schedules

## DevOps Lifecycle

### CONTINUOUS CYCLE

#### 1. PLAN

- Define features
- Create user stories
- Sprint planning

#### 2. CODE

- Write code
- Version control (Git)
- Code review

#### 3. BUILD

- Compile code
- Run unit tests
- Create artifacts

#### 4. TEST

- Integration tests
- Security tests
- Performance tests

#### 5. RELEASE

- Staging deployment
- User acceptance testing
- Approval gates

#### 6. DEPLOY

- Production deployment
- Blue-green deployment
- Canary releases

#### 7. OPERATE

- Monitor applications
- Manage infrastructure
- Handle incidents

#### 8. MONITOR

- Collect metrics

Analyze logs  
User feedback

→ Back to PLAN (Continuous Improvement)

## DevOps Tools Ecosystem

### Version Control

```
# Git for source code
git init
git add .
git commit -m "New feature"
git push origin main
```

### Continuous Integration/Continuous Deployment (CI/CD)

```
# Jenkins, GitLab CI, GitHub Actions
name: CI/CD Pipeline
on: [push]
jobs:
  build:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v2
      - name: Build
        run: docker build -t myapp .
      - name: Test
        run: docker run myapp npm test
      - name: Deploy
        run: kubectl apply -f deployment.yaml
```

### Containerization

```
# Docker
docker build -t myapp .
docker run -d -p 80:80 myapp
```

### Orchestration

```
# Kubernetes
kubectl create deployment myapp --image=myapp:latest
kubectl scale deployment myapp --replicas=10
```

### Infrastructure as Code

```
# Terraform
resource "aws_instance" "web" {
```

```
    ami           = "ami-0c55b159cbfafe1f0"
    instance_type = "t2.micro"
}
```

## Monitoring & Logging

```
# Prometheus, Grafana, ELK Stack
# Monitor CPU, memory, requests, errors
# Alert on anomalies
```

## Docker's Role in DevOps

Docker Solves Multiple DevOps Challenges:

### 1. Environment Consistency

```
# Define environment once
FROM node:18
# Works everywhere: dev, test, prod
```

### 2. Fast Deployment

```
# Old way: Hours or days
# Docker way: Seconds
docker run -d myapp:latest
```

### 3. Easy Scaling

```
# Scale from 1 to 100 instances
docker-compose up --scale web=100
```

### 4. Microservices Architecture

```
# docker-compose.yml
services:
  frontend:
    image: frontend:latest
  backend:
    image: backend:latest
  database:
    image: postgres:14
  cache:
    image: redis:7
```

### 5. Infrastructure as Code

```
# Dockerfile IS your infrastructure
FROM python:3.10
```

```
RUN apt-get update  
RUN pip install django  
# Version controlled, reproducible
```

## DevOps Metrics (DORA Metrics)

### Elite Performers:

Deployment Frequency:	Multiple times per day
Lead Time for Changes:	Less than 1 hour
Time to Restore Service:	Less than 1 hour
Change Failure Rate:	0-15%

### Low Performers:

Deployment Frequency:	Once per month or less
Lead Time for Changes:	1-6 months
Time to Restore Service:	1 week to 1 month
Change Failure Rate:	46-60%

### Docker Helps You Become Elite:

```
# Deploy in seconds  
docker push myapp:v2.1.5  
docker pull myapp:v2.1.5  
docker run myapp:v2.1.5  
  
# Rollback in seconds if needed  
docker run myapp:v2.1.4
```

## Real-World DevOps Success Story

### Netflix Example:

#### Before DevOps:

- Quarterly releases
- Weeks of testing
- Frequent outages
- Slow recovery

#### After DevOps + Docker + Microservices:

- Thousands of deployments per day
  - Minutes of testing per service
  - Minimal downtime
  - Auto-recovery
  - Serves 200+ million subscribers
-

## Key Takeaways

1. **Docker** is a containerization platform that packages apps with their dependencies
  2. **Containers** are lightweight, isolated, and portable environments
  3. **DevOps** is a culture that brings development and operations together
  4. **Docker + DevOps** = Fast, reliable, scalable deployments
  5. **History** shows containers evolved over decades, Docker made them accessible
- 

## What's Next?

In the next sections, we'll cover: - Docker Architecture - Installing Docker - Working with Containers - Building Docker Images - And much more!

**Remember:** Docker is not just a tool, it's a game-changer for modern software development!